

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [42]: kisan=pd.read_csv('data_folder/kisan_net_log.csv')
kisan.head()
```


C:\Users\Atul kumar\anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3063: DtypeWarning: Columns (19,29,30) have mixed types.Specify dtype option on import or set low\_memory=False.

interactivity=interactivity, compiler=compiler, result=result)

Out[42]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ow
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	Mendozaberg	OK 22690"	NaN	NaN	NaN	NaN	NaN	NaN	
2	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MOR
3	Loganmouth	SD 05113"	NaN	NaN	NaN	NaN	NaN	NaN	
4	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	

5 rows × 31 columns




```
In [43]: kisan1=kisan.iloc[0::2,0:27]
kisan1.head()
```

Out[43]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ow
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
2	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MOR
4	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
6	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
8	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MOR

5 rows × 27 columns



EDA

Target Variable-Loan\_Status

```
In [44]: kisan1['loan_status'].value_counts()
```

```
Out[44]: Fully Paid      78498
Charged Off    19294
Jun-13         125
Aug-13         107
May-13          99
...
35000          1
40000          1
Sep-07         1
42000          1
81090          1
Name: loan_status, Length: 127, dtype: int64
```

```
In [45]: kisan2=kisan1[(kisan1['loan_status']=='Fully Paid')|(kisan1['loan_status']=='Charged Off')]
kisan2.shape
```

```
Out[45]: (97792, 27)
```

```
In [46]: kisan2['loan_status'].value_counts(normalize=True)
#Highly Imbalanced dataset
```

```
Out[46]: Fully Paid      0.802704
Charged Off    0.197296
Name: loan_status, dtype: float64
```

```
In [47]: kisan2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 97792 entries, 0 to 200488
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             97792 non-null  object
1   term                  97792 non-null  object
2   int_rate              97792 non-null  float64
3   installment           97792 non-null  float64
4   grade                 97792 non-null  object
5   sub_grade             97792 non-null  object
6   emp_title             92061 non-null  object
7   emp_length            93190 non-null  object
8   home_ownership        97792 non-null  object
9   annual_inc            97792 non-null  object
10  verification_status   97792 non-null  object
11  issue_d               97792 non-null  object
12  loan_status           97792 non-null  object
13  purpose               97792 non-null  object
14  title                 97351 non-null  object
15  dti                   97792 non-null  object
16  earliest_cr_line      97792 non-null  object
17  open_acc              97792 non-null  object
18  pub_rec               97792 non-null  object
19  revol_bal             97792 non-null  object
20  revol_util            97723 non-null  float64
21  total_acc             97792 non-null  float64
22  initial_list_status   97792 non-null  object
23  application_type      97792 non-null  object
24  mort_acc              88924 non-null  object
25  pub_rec_bankruptcies  97607 non-null  object
26  address               97782 non-null  object
dtypes: float64(4), object(23)
memory usage: 20.9+ MB
```

In [48]: kisan2.isnull().sum()

```
Out[48]: loan_amnt      0
term      0
int_rate  0
installment  0
grade     0
sub_grade  0
emp_title  5731
emp_length 4602
home_ownership  0
annual_inc  0
verification_status  0
issue_d      0
loan_status  0
purpose     0
title      441
dti         0
earliest_cr_line  0
open_acc    0
pub_rec     0
revol_bal   0
revol_util  69
total_acc   0
initial_list_status  0
application_type  0
mort_acc    8868
pub_rec_bankruptcies  185
address     10
dtype: int64
```

```
In [49]: import re
pattern='[a-zA-Z]'
cols=['loan_amnt', 'annual_inc', 'dti', 'revol_bal', 'mort_acc']
for i in cols:
    kisan2[i]=list(map(lambda x: 'nan' if re.search(pattern, str(x)) else x, kisan2[i]))
    kisan2=kisan2[kisan2[i]!='nan']
kisan2.shape
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) (https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy)

Out[49]: (88741, 27)

Changing the data types of some features to floats

```
In [50]: cols=['loan_amnt','int_rate','installment','annual_inc','dti',
            'revol_bal','revol_util']
for i in cols:
    kisan2[i]=list(map(lambda x:float(x),kisan2[i]))
kisan2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 88741 entries, 0 to 200488
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            88741 non-null  float64
1   term                                88741 non-null  object
2   int_rate                            88741 non-null  float64
3   installment                          88741 non-null  float64
4   grade                               88741 non-null  object
5   sub_grade                           88741 non-null  object
6   emp_title                           83607 non-null  object
7   emp_length                          84403 non-null  object
8   home_ownership                      88741 non-null  object
9   annual_inc                          88741 non-null  float64
10  verification_status                 88741 non-null  object
11  issue_d                             88741 non-null  object
12  loan_status                         88741 non-null  object
13  purpose                             88741 non-null  object
14  title                               88302 non-null  object
15  dti                                 88741 non-null  float64
16  earliest_cr_line                    88741 non-null  object
17  open_acc                            88741 non-null  object
18  pub_rec                             88741 non-null  object
19  revol_bal                           88741 non-null  float64
20  revol_util                          88684 non-null  float64
21  total_acc                           88741 non-null  float64
22  initial_list_status                 88741 non-null  object
23  application_type                    88741 non-null  object
24  mort_acc                            88741 non-null  object
25  pub_rec_bankruptcies                88741 non-null  object
26  address                             88741 non-null  object
dtypes: float64(8), object(19)
memory usage: 19.0+ MB
```

```
In [51]: cols=['open_acc', 'pub_rec', 'mort_acc',
              'pub_rec_bankruptcies']
for i in cols:
    kisan2[i]=list(map(lambda x:int(x),kisan2[i]))
kisan2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 88741 entries, 0 to 200488
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            88741 non-null  float64
1   term                                 88741 non-null  object
2   int_rate                             88741 non-null  float64
3   installment                          88741 non-null  float64
4   grade                                88741 non-null  object
5   sub_grade                            88741 non-null  object
6   emp_title                            83607 non-null  object
7   emp_length                           84403 non-null  object
8   home_ownership                       88741 non-null  object
9   annual_inc                           88741 non-null  float64
10  verification_status                  88741 non-null  object
11  issue_d                              88741 non-null  object
12  loan_status                          88741 non-null  object
13  purpose                              88741 non-null  object
14  title                                88302 non-null  object
15  dti                                  88741 non-null  float64
16  earliest_cr_line                     88741 non-null  object
17  open_acc                             88741 non-null  int64
18  pub_rec                              88741 non-null  int64
19  revol_bal                            88741 non-null  float64
20  revol_util                           88684 non-null  float64
21  total_acc                            88741 non-null  float64
22  initial_list_status                  88741 non-null  object
23  application_type                     88741 non-null  object
24  mort_acc                             88741 non-null  int64
25  pub_rec_bankruptcies                 88741 non-null  int64
26  address                              88741 non-null  object
dtypes: float64(8), int64(4), object(15)
memory usage: 19.0+ MB
```

```
In [52]: kisan2['title'].value_counts()
```

```
Out[52]: Debt consolidation                37962
Credit card refinancing                  12970
Home improvement                          3865
Other                                    3179
Debt Consolidation                        2191
...
Home Improvement/Consolidation            1
loan payback then purchase new shop       1
consolidation of credit cards             1
Debt Consolidation & CC Reduction         1
Persona                                  1
Name: title, Length: 10289, dtype: int64
```

```
In [53]: kisan2['purpose'].value_counts()
```

```
Out[53]: debt_consolidation    53511  
credit_card    19352  
home_improvement    5359  
other    4275  
major_purchase    1723  
small_business    939  
medical    889  
car    808  
moving    599  
vacation    549  
house    468  
wedding    215  
renewable_energy    53  
educational    1  
Name: purpose, dtype: int64
```

Categorising purpose into:debt\_consolidation,credit\_card and others

```
In [54]: z=['debt_consolidation','credit_card']  
kisan2['purpose']=list(map(lambda x:x if x in z else 'others',kisan2['purpose']))  
kisan2['purpose'].value_counts()
```

```
Out[54]: debt_consolidation    53511  
credit_card    19352  
others    15878  
Name: purpose, dtype: int64
```

```
In [55]: kisan2['home_ownership'].value_counts()
```

```
Out[55]: MORTGAGE    45023  
RENT    35083  
OWN    8613  
OTHER    13  
NONE    7  
ANY    2  
Name: home_ownership, dtype: int64
```

Categorising home\_ownership into:Mortgage,Rent and Others

```
In [56]: z=['MORTGAGE','RENT']  
kisan2['home_ownership']=list(map(lambda x:x if x in z else 'OTHERS',kisan2['home_ownership']))  
kisan2['home_ownership'].value_counts()
```

```
Out[56]: MORTGAGE    45023  
RENT    35083  
OTHERS    8635  
Name: home_ownership, dtype: int64
```

```
In [57]: kisan2['verification_status'].value_counts()
```

```
Out[57]: Verified          31144  
Source Verified    30265  
Not Verified       27332  
Name: verification_status, dtype: int64
```

Merging the Veified and Source Verified for verification\_status

```
In [58]: z=['Verified','Source Verified']  
kisan2['verification_status']=list(map(lambda x:'Verified' if x in z else x,kisan2['verification_status'].value_counts()))
```

```
Out[58]: Verified          61409  
Not Verified       27332  
Name: verification_status, dtype: int64
```

pub\_rec:No. of derogatory public records

```
In [59]: kisan2['pub_rec'].value_counts()
```

```
Out[59]: 0      74769  
1      11950  
2       1393  
3        360  
4        147  
5         63  
6         27  
8         15  
7          8  
10         3  
11         2  
9          2  
19         1  
13         1  
Name: pub_rec, dtype: int64
```

```
In [60]: z=[0,1]  
kisan2['pub_rec']=list(map(lambda x:x if x in z else 1,kisan2['pub_rec']))  
kisan2['pub_rec'].value_counts()
```

```
Out[60]: 0      74769  
1      13972  
Name: pub_rec, dtype: int64
```



```
In [61]: kisan2['term'].value_counts()
```

```
Out[61]: 36 months    67798
        60 months    20943
        Name: term, dtype: int64
```

```
In [62]: kisan2['is_36mnths']=list(map(lambda x:1 if x==' 36 months'else 0,kisan2['term']))
        kisan2['is_36mnths'].value_counts()
```

```
Out[62]: 1    67798
        0    20943
        Name: is_36mnths, dtype: int64
```

```
In [63]: kisan2['is_verified']=list(map(lambda x:1 if x=='Verified'else 0,kisan2['verified']))
        kisan2['is_verified'].value_counts()
```

```
Out[63]: 1    61409
        0    27332
        Name: is_verified, dtype: int64
```

```
In [64]: kisan2['grade'].value_counts()
```

```
Out[64]: B    25968
        C    24431
        D    14598
        A    13215
        E     7207
        F     2642
        G      680
        Name: grade, dtype: int64
```

```
In [65]: d={'A':0, 'B':1, 'C':2, 'D':3, 'E':4, 'F':5, 'G':6}
kisan2['mod_grade']=list(map(lambda x:d[x],kisan2['grade']))
kisan2[['grade', 'mod_grade']].head(20)
```

Out[65]:

	grade	mod_grade
0	B	1
2	B	1
4	B	1
6	A	0
8	C	2
10	C	2
12	A	0
14	B	1
16	B	1
18	C	2
20	B	1
22	C	2
24	B	1
26	C	2
28	A	0
30	A	0
32	E	4
34	C	2
36	A	0
38	A	0

```
In [66]: kisan2=pd.get_dummies(kisan2,columns=['home_ownership','purpose'],drop_first=True)
```

```
In [67]: cols=['home_ownership_OTHERS', 'home_ownership_RENT', 'purpose_debt_consolidation',
              'purpose_others']
for i in cols:
    kisan2[i]=list(map(lambda x:int(x),kisan2[i]))
kisan2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 88741 entries, 0 to 200488
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            88741 non-null  float64
1   term                                88741 non-null  object
2   int_rate                            88741 non-null  float64
3   installment                         88741 non-null  float64
4   grade                               88741 non-null  object
5   sub_grade                           88741 non-null  object
6   emp_title                           83607 non-null  object
7   emp_length                         84403 non-null  object
8   annual_inc                         88741 non-null  float64
9   verification_status                88741 non-null  object
10  issue_d                             88741 non-null  object
11  loan_status                        88741 non-null  object
12  title                              88302 non-null  object
13  dti                                88741 non-null  float64
14  earliest_cr_line                   88741 non-null  object
15  open_acc                          88741 non-null  int64
16  pub_rec                           88741 non-null  int64
17  revol_bal                         88741 non-null  float64
18  revol_util                        88684 non-null  float64
19  total_acc                         88741 non-null  float64
20  initial_list_status                88741 non-null  object
21  application_type                   88741 non-null  object
22  mort_acc                          88741 non-null  int64
23  pub_rec_bankruptcies               88741 non-null  int64
24  address                            88741 non-null  object
25  is_36mnths                        88741 non-null  int64
26  is_verified                       88741 non-null  int64
27  mod_grade                         88741 non-null  int64
28  home_ownership_OTHERS              88741 non-null  int64
29  home_ownership_RENT                88741 non-null  int64
30  purpose_debt_consolidation         88741 non-null  int64
31  purpose_others                     88741 non-null  int64
dtypes: float64(8), int64(11), object(13)
memory usage: 22.3+ MB
```

Type *Markdown* and LaTeX:  $\alpha^2$

```
In [68]: kisan2['application_type'].value_counts(normalize=True)
```

```
Out[68]: INDIVIDUAL    0.998073
         JOINT         0.001127
         DIRECT_PAY    0.000800
         Name: application_type, dtype: float64
```

```
In [69]: kisan2['initial_list_status'].value_counts(normalize=True)
```

```
Out[69]: f    0.555775
         w    0.444225
         Name: initial_list_status, dtype: float64
```

```
In [70]: d={'f':0,'w':1}
         kisan2['initial_list_status']=list(map(lambda x:d[x],kisan2['initial_list_status']))
         kisan2['initial_list_status'].value_counts()
```

```
Out[70]: 0    49320
         1    39421
         Name: initial_list_status, dtype: int64
```

```
In [72]: kisan2['pub_rec_bankruptcies'].value_counts()
```

```
Out[72]: 0    77846
         1    10303
         2     475
         3     79
         4     25
         5      8
         6      4
         7      1
         Name: pub_rec_bankruptcies, dtype: int64
```

pub\_rec\_bankruptcies: Number of public record bankruptcies

```
In [73]: z=[0,1]
         kisan2['pub_rec_bankruptcies']=list(map(lambda x:x if x in z else 1,kisan2['pub_rec_bankruptcies']))
         kisan2['pub_rec_bankruptcies'].value_counts()
```

```
Out[73]: 0    77846
         1    10895
         Name: pub_rec_bankruptcies, dtype: int64
```

```
In [86]: kisan3=kisan2.dropna()
         kisan3.shape
```

```
Out[86]: (83135, 32)
```

```
In [87]: kisan3.isnull().sum()
```

```
Out[87]: loan_amnt      0
term      0
int_rate  0
installment  0
grade     0
sub_grade 0
emp_title  0
emp_length 0
annual_inc 0
verification_status 0
issue_d    0
loan_status 0
title     0
dti       0
earliest_cr_line 0
open_acc  0
pub_rec   0
revol_bal 0
revol_util 0
total_acc 0
initial_list_status 0
application_type 0
mort_acc  0
pub_rec_bankruptcies 0
address    0
is_36mnths 0
is_verified 0
mod_grade  0
home_ownership_OTHERS 0
home_ownership_RENT 0
purpose_debt_consolidation 0
purpose_others 0
dtype: int64
```

Splitting the Kisan3 dataset into train,validation and test sets

```
In [89]: XS=kisan3[['loan_amnt','int_rate', 'installment','mod_grade','sub_grade',
                    'annual_inc','is_verified','dti','open_acc', 'pub_rec',
                    'revol_bal', 'revol_util', 'total_acc','initial_list_status',
                    'mort_acc','pub_rec_bankruptcies','is_36mnths',
                    'home_ownership_OTHERS', 'home_ownership_RENT',
                    'purpose_debt_consolidation', 'purpose_others']]
```

```
In [91]: XS.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 83135 entries, 0 to 200488
Data columns (total 21 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   loan_amnt                             83135 non-null  float64
 1   int_rate                             83135 non-null  float64
 2   installment                           83135 non-null  float64
 3   mod_grade                             83135 non-null  int64   
 4   sub_grade                             83135 non-null  object  
 5   annual_inc                             83135 non-null  float64
 6   is_verified                           83135 non-null  int64   
 7   dti                                    83135 non-null  float64
 8   open_acc                              83135 non-null  int64   
 9   pub_rec                               83135 non-null  int64   
10  revol_bal                             83135 non-null  float64
11  revol_util                             83135 non-null  float64
12  total_acc                             83135 non-null  float64
13  initial_list_status                   83135 non-null  int64   
14  mort_acc                              83135 non-null  int64   
15  pub_rec_bankruptcies                  83135 non-null  int64   
16  is_36mnths                           83135 non-null  int64   
17  home_ownership_OTHERS                 83135 non-null  int64   
18  home_ownership_RENT                   83135 non-null  int64   
19  purpose_debt_consolidation            83135 non-null  int64   
20  purpose_others                        83135 non-null  int64   
dtypes: float64(8), int64(12), object(1)
memory usage: 14.0+ MB
```

```
In [92]: y=kisan3['loan_status']
```

```
In [96]: kisan3['y_obs']=list(map(lambda x:0 if x=='Fully Paid' else 1,kisan3['loan_status']
kisan3[['loan_status','y_obs']].head(14)
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

"""Entry point for launching an IPython kernel.

Out[96]:

	loan_status	y_obs
0	Fully Paid	0
2	Fully Paid	0
4	Fully Paid	0
6	Fully Paid	0
8	Charged Off	1
10	Fully Paid	0
12	Fully Paid	0
14	Fully Paid	0
16	Fully Paid	0
18	Fully Paid	0
20	Fully Paid	0
22	Fully Paid	0
24	Fully Paid	0
26	Fully Paid	0

```
In [97]: y=kisan3['y_obs']
```

```
In [98]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(XS,y,test_size=0.15,random_state=2)
```

```
In [99]: X_trainn,X_dev,y_trainn,y_dev=train_test_split(X_train,y_train,test_size=0.15,ran
```

```
In [100]: print(X_trainn.shape)
          print(X_dev.shape)
          print(X_test.shape)
          print(y_trainn.shape)
          print(y_dev.shape)
          print(y_test.shape)
```

```
(60064, 21)
(10600, 21)
(12471, 21)
(60064,)
(10600,)
(12471,)
```

```
In [101]: X_trainn=pd.concat([X_trainn,y_trainn],axis=1)
```

```
In [103]: z=X_trainn.groupby('sub_grade')['y_obs'].mean().to_dict()
          X_trainn['sub_grade']=list(map(lambda x:z[x],X_trainn['sub_grade']))
          X_trainn['sub_grade']
```

```
Out[103]: 14488      0.192450
          94254      0.154802
          24670      0.341498
          20048      0.168598
          55286      0.075792
          ...
          71550      0.108300
          86534      0.347036
          117018     0.085448
          190226     0.213446
          184816     0.125066
          Name: sub_grade, Length: 60064, dtype: float64
```



```
In [104]: X_dev['sub_grade']=list(map(lambda x:z[x],X_dev['sub_grade']))
X_dev['sub_grade']
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

"""Entry point for launching an IPython kernel.

```
Out[104]: 39684      0.404795
16006      0.122754
165846     0.154802
151720     0.062593
15474      0.192450
...
192714     0.154802
38598      0.122754
132708     0.028833
90502      0.241294
179298     0.322265
Name: sub_grade, Length: 10600, dtype: float64
```

```
In [105]: X_test['sub_grade']=list(map(lambda x:z[x],X_test['sub_grade']))
X_test['sub_grade']
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

"""Entry point for launching an IPython kernel.

```
Out[105]: 56938      0.322265
63494      0.154802
120154     0.240930
26788      0.265314
77820      0.108300
...
143328     0.125066
177034     0.265314
48176      0.240930
49746      0.101064
110436     0.154802
Name: sub_grade, Length: 12471, dtype: float64
```

```
In [107]: cols=['loan_amnt','int_rate','installment','annual_inc','dti',  
              'open_acc','revol_bal','revol_util','total_acc','mort_acc',  
              'mod_grade']  
d={}  
for i in cols:  
    d[i]=[X_trainn[i].mean(),X_trainn[i].std()]  
print(d)
```


```
{'loan_amnt': [14491.336324587108, 8388.981997556977], 'int_rate': [13.78460392  
2482301, 4.490693801881828], 'installment': [444.271407332167, 251.863757655670  
16], 'annual_inc': [75394.66977274243, 48671.130821498744], 'dti': [17.68224127  
597239, 8.108662503973125], 'open_acc': [11.59971030900373, 5.213968658647141],  
'revol_bal': [16033.318560202451, 19607.84551585026], 'revol_util': [54.1856752  
7970166, 23.89877657577768], 'total_acc': [25.88725359616409, 11.92963296750360  
4], 'mort_acc': [1.8127164358018113, 2.148206675715537], 'mod_grade': [1.852124  
400639318, 1.321668894368757]}
```

```
In [108]: cols=['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
               'open_acc', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
               'mod_grade']
for i in cols:
    mean=X_trainn[i].mean()
    std=X_trainn[i].std()
    X_trainn[i]=(X_trainn[i]-mean)/std
X_trainn.head(20)
```

Out[108]:

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	is_verified	d
14488	-0.869156	0.339679	-0.768596	0.111886	0.192450	-0.735439	0	0.55468
94254	0.418247	0.068006	0.681792	-0.644734	0.154802	0.053529	0	-0.05330
24670	0.060635	0.907520	-0.256136	0.868505	0.341498	-0.110839	1	0.43012
20048	-0.296977	-0.176945	-0.158822	0.111886	0.168598	-0.234116	0	0.85436
55286	1.610286	-1.310400	1.714652	-1.401353	0.075792	1.347931	1	-0.07057
138636	0.418247	0.442113	0.740554	0.868505	0.265314	-0.521760	1	0.63114
47806	-1.155246	-1.595434	-1.178778	-1.401353	0.040411	-0.850497	0	-1.59363
151938	-1.369813	-1.726816	-1.401398	-1.401353	0.028833	-0.727221	0	0.99372
153512	0.626854	-0.150223	0.882376	-0.644734	0.125066	-0.460122	1	-0.36901
151486	-1.012201	-1.582073	-1.031833	-1.401353	0.062593	0.505543	0	-0.58483
106528	2.444714	1.381835	1.917023	1.625124	0.391566	0.711003	1	0.50535
167006	-0.589027	0.713786	-0.412252	0.868505	0.298160	-0.932682	1	1.51415
129208	0.418247	0.994367	0.060622	0.868505	0.268987	-0.460122	1	0.56825
162968	-0.907898	1.604517	-0.735641	1.625124	0.391566	-1.035412	1	1.06401
163026	-0.726112	-0.319462	-0.650595	-0.644734	0.125066	0.074075	1	-0.53674
77624	1.610286	1.381835	1.180831	0.868505	0.322265	0.197352	1	0.20074
17694	-0.654589	0.402031	-0.514847	0.111886	0.241294	-0.809405	1	1.41302
20454	0.954665	-0.266018	1.228476	0.111886	0.192450	-0.521760	1	-1.61829
23510	-0.535385	-0.522103	-0.455768	-0.644734	0.125066	0.505543	1	-1.25695
161670	-0.535385	-0.176945	-0.426307	0.111886	0.168598	-0.829951	1	1.53635

20 rows × 22 columns



```
In [109]: cols=['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
              'open_acc', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
              'mod_grade']
for i in cols:
    X_dev[i]=(X_dev[i]-d[i][0])/d[i][1]
print(X_dev)
```

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	\
39684	2.444714	1.929634	2.109746	1.625124	0.404795	7.285743	
16006	-1.310211	-0.733651	-1.312302	-0.644734	0.122754	-0.398484	
165846	0.060635	-0.399627	0.213920	-0.644734	0.154802	-0.213570	
151720	1.610286	-1.372751	1.700318	-1.401353	0.062593	1.841037	
15474	-0.177773	-0.176945	-0.589769	0.111886	0.192450	0.197352	
...	...	...	...	...	...	...	
192714	-0.296977	-0.502061	-0.715353	-0.644734	0.154802	-0.727221	
38598	-0.714191	-0.399627	-0.643131	-0.644734	0.122754	-0.419030	
132708	-0.535385	-1.884921	-0.568249	-1.401353	0.028833	-0.012218	
90502	1.133471	0.667023	0.593013	0.111886	0.241294	0.300082	
179298	1.783132	0.842942	1.178012	0.868505	0.322265	-0.336846	

	is_verified	dti	open_acc	pub_rec	...	revol_util	total_acc
39684	1	-0.561405	0.268565	0	...	1.172207	1.434474
16006	1	-0.758725	0.652150	0	...	-1.601993	-0.912623
165846	0	0.622514	0.076773	0	...	-0.288118	0.512400
151720	1	-0.996742	1.035735	0	...	-0.555914	1.182999
15474	0	0.733750	1.073000	0	...	0.757000	0.661110

```
In [110]: cols=['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
              'open_acc', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
              'mod_grade']
for i in cols:
    X_test[i]=(X_test[i]-d[i][0])/d[i][1]
print(X_test)
```

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	\
56938	1.875515	1.210369	2.648371	0.868505	0.322265	-0.069747	
63494	1.133471	-0.502061	1.379709	-0.644734	0.154802	-0.143230	
120154	0.179839	0.339679	-0.242279	0.111886	0.240930	-0.316300	
26788	-0.392340	0.406484	-0.691729	0.868505	0.265314	-0.710784	
77820	0.060635	-0.588908	0.189819	-0.644734	0.108300	0.094621	
...	...	...	...	...	...	...	
143328	-0.535385	-0.176945	-0.426307	-0.644734	0.125066	-0.521760	
177034	-0.296977	0.406484	-0.615140	0.868505	0.265314	-0.768313	
48176	-0.773793	0.045738	-0.678468	0.111886	0.240930	-0.747767	
49746	-0.535385	-0.844993	-0.482965	-0.644734	0.101064	1.121924	
110436	-0.010888	0.068006	0.192638	-0.644734	0.154802	-0.656337	

	is_verified	dti	open_acc	pub_rec	...	revol_util	total_acc
56938	1	0.951792	-0.882190	0	...	1.327027	-0.912623
63494	1	0.967824	1.227527	0	...	-1.003636	0.847700
120154	1	-1.043605	-0.690397	0	...	0.732854	-0.158199
26788	1	1.932225	0.268565	0	...	0.448321	-0.912623
77820	1	1.336566	0.306010	0	...	1.456711	0.800151

```
In [111]: X_trainn=X_trainn.drop('y_obs',axis=1)
X_trainn.columns
```

```
Out[111]: Index(['loan_amnt', 'int_rate', 'installment', 'mod_grade', 'sub_grade',
               'annual_inc', 'is_verified', 'dti', 'open_acc', 'pub_rec', 'revol_bal',
               'revol_util', 'total_acc', 'initial_list_status', 'mort_acc',
               'pub_rec_bankruptcies', 'is_36mnths', 'home_ownership_OTHERS',
               'home_ownership_RENT', 'purpose_debt_consolidation', 'purpose_others'],
              dtype='object')
```

```
In [113]: from sklearn.linear_model import LogisticRegression
```

```
In [118]: y_trainn.value_counts(normalize=True)
```

```
Out[118]: 0    0.801345
          1    0.198655
          Name: y_obs, dtype: float64
```

```
In [120]: from sklearn import metrics
```

Making a Logistic regression model

```

In [121]: train_cost=[]
          val_cost=[]
          lambdaa=[]
          for i in np.arange(0.01,100,2):
              lambdaa.append(i)
              model=LogisticRegression(C=1/i,class_weight={0:0.1,1:0.6})
              model.fit(X_trainn,y_trainn)
              y_train_probs=model.predict_proba(X_trainn)
              y_val_probs=model.predict_proba(X_dev)
              train_loss=metrics.log_loss(y_trainn,y_train_probs)
              val_loss=metrics.log_loss(y_dev,y_val_probs)
              train_cost.append(train_loss)
              val_cost.append(val_loss)
          plt.plot(lambdaa,train_cost,label='train')
          plt.plot(lambdaa,val_cost,label='val')
          plt.legend()
          plt.show()

```

C:\Users\Atul kumar\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)



```

In [122]: z=np.argmin(val_cost)
          z

```

Out[122]: 0

```
In [123]: model=LogisticRegression(C=1,class_weight={0:0.1,1:0.6})
model.fit(X_trainn,y_trainn)
print(model.score(X_trainn,y_trainn))
print(model.score(X_dev,y_dev))
print(model.score(X_test,y_test))
```

```
0.5457844965370272
0.5323584905660378
0.5452650148344158
```

```
In [124]: from sklearn.metrics import confusion_matrix, precision_score, recall_score
ypred=model.predict(X_test)
confusion_matrix(y_test,ypred)
```

```
Out[124]: array([[4788, 5176],
                [ 495, 2012]], dtype=int64)
```

```
In [125]: print(precision_score(y_test,ypred))
print(recall_score(y_test,ypred))
```

```
0.27991096271563715
0.8025528520143598
```

Removing the outliers which are on the wrong side of the plane  $\theta_T^T x + \theta_{not}$

```
In [126]: z=np.array(list(model.intercept_)+list(model.coef_[0]))
z=z.T
z
```

```
Out[126]: array([-0.37986903, -0.09505128, -0.24157436,  0.16026868,  0.29393671,
                 3.91784125, -0.07684596,  0.10735067,  0.21085049,  0.10616443,
                 0.08805733, -0.07201472,  0.08561362, -0.1274847 , -0.0507649 ,
                -0.07308559, -0.08033209, -0.49011554,  0.10939622,  0.26325411,
                 0.02387176,  0.12101527])
```

```
In [127]: X_trainn.insert(0,'x_not',1)
```

In [128]: X\_trainn

Out[128]:

	x_not	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	is_verified
14488	1	-0.869156	0.339679	-0.768596	0.111886	0.192450	-0.735439	0
94254	1	0.418247	0.068006	0.681792	-0.644734	0.154802	0.053529	0
24670	1	0.060635	0.907520	-0.256136	0.868505	0.341498	-0.110839	1
20048	1	-0.296977	-0.176945	-0.158822	0.111886	0.168598	-0.234116	0
55286	1	1.610286	-1.310400	1.714652	-1.401353	0.075792	1.347931	1
...	...	...	...	...	...	...	...	...
71550	1	0.656655	-0.956334	0.779384	-0.644734	0.108300	-0.117003	1
86534	1	0.299043	1.159152	-0.013386	1.625124	0.347036	-0.291645	1
117018	1	0.662615	-1.087717	0.763860	-1.401353	0.085448	-0.624491	1
190226	1	-1.369813	0.266194	-1.351133	0.111886	0.213446	0.197352	0
184816	1	-1.143325	-0.622310	-1.127083	-0.644734	0.125066	-0.172477	1

60064 rows × 22 columns

In [129]: z2=X\_trainn.dot(z)  
z3=1/(1+np.exp(-z2))  
z3

Out[129]: 14488 0.555917  
94254 0.504452  
24670 0.769651  
20048 0.539214  
55286 0.339882  
...  
71550 0.491943  
86534 0.854377  
117018 0.431970  
190226 0.471579  
184816 0.391628  
Length: 60064, dtype: float64

In [130]: X\_trainn['log\_loss']=-1\*(y\_trainn\*np.log(z3)+(1-y\_trainn)\*(1-np.log(z3)))



```
In [131]: X_trainn['log_loss'].describe()
```

```
Out[131]: count    60064.000000
          mean      -1.268352
          std       0.921445
          min      -5.582546
          25%      -1.851313
          50%      -1.563946
          75%      -1.232072
          max       2.381357
          Name: log_loss, dtype: float64
```

```
In [132]: X_trainn=pd.concat([X_trainn,y_trainn],axis=1)
```

```
In [133]: q05=X_trainn['log_loss'].quantile(0.05)
          q95=X_trainn['log_loss'].quantile(0.95)
          X_trainn=X_trainn[(X_trainn['log_loss']>q05)&(X_trainn['log_loss']<q95)]
```

```
In [134]: y_trainn=X_trainn['y_obs']
          y_trainn
```

```
Out[134]: 14488      0
          94254      0
          24670      0
          20048      0
          55286      0
          ..
          71550      0
          86534      1
          117018     0
          190226     0
          184816     0
          Name: y_obs, Length: 54056, dtype: int64
```

```
In [136]: X_trainn=X_trainn.drop(['x_not','y_obs','log_loss'],axis=1)
```

Performing LogisticRegression after removing outliers

```

In [173]: train_cost=[]
          val_cost=[]
          lambdaa=[]
          for i in np.arange(0.1,1000,5):
              lambdaa.append(i)
              model=LogisticRegression(C=1/i,class_weight={0:0.1,1:0.9})
              model.fit(X_trainn,y_trainn)
              y_train_probs=model.predict_proba(X_trainn)
              y_val_probs=model.predict_proba(X_dev)
              train_loss=metrics.log_loss(y_trainn,y_train_probs)
              val_loss=metrics.log_loss(y_dev,y_val_probs)
              train_cost.append(train_loss)
              val_cost.append(val_loss)
          plt.plot(lambdaa,train_cost,label='train')
          plt.plot(lambdaa,val_cost,label='val')
          plt.legend()
          plt.show()

```

C:\Users\Atul kumar\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)



In [ ]:

```
In [174]: z=np.argmin(val_cost)
lbest=0.1+z*5
lbest
```

Out[174]: 785.1

```
In [175]: model=LogisticRegression(C=1/lbest,class_weight={0:0.1,1:0.9})
model.fit(X_trainn,y_trainn)
print(model.score(X_trainn,y_trainn))
print(model.score(X_dev,y_dev))
print(model.score(X_test,y_test))
```

```
0.6240565339647772
0.6033962264150944
0.6118194210568519
```

```
In [178]: ypred=model.predict(X_test)
confusion_matrix(y_test,ypred)
```

Out[178]: array([[5850, 4114],  
[ 727, 1780]], dtype=int64)

```
In [177]: print(precision_score(y_test,ypred))
print(recall_score(y_test,ypred))
```

```
0.30200203596878183
0.7100119664938173
```

Introducing polynomial terms in the logistic regression model

```
In [180]: cols=['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
              'open_acc', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
              'mod_grade']
t=0
for i in cols:
    z='f'+str(t)
    XS[z]=XS[i]**2
    t+=1
print(XS)
```

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	\
0	10000.0	11.44	329.48	1	B4	117000.0	
2	8000.0	11.99	265.68	1	B5	65000.0	
4	15600.0	10.49	506.97	1	B3	43057.0	
6	7200.0	6.49	220.65	0	A2	54000.0	
8	24375.0	17.27	609.33	2	C5	55000.0	
...	...	...	...	...	...	...	
200476	30225.0	17.27	755.57	2	C5	68000.0	
200478	6000.0	5.32	180.69	0	A1	48000.0	
200480	15450.0	7.90	483.44	0	A4	42000.0	
200486	5000.0	13.53	169.75	1	B5	43000.0	
200488	15000.0	14.30	514.86	2	C1	44000.0	

	is_verified	dti	open_acc	pub_rec	...	f1	f2	\
0	0	26.24	16	0	...	130.8736	108557.0704	
2	0	22.05	17	0	...	143.7601	70585.8624	
4	1	12.79	13	0	...	110.0401	257018.5809	
6	0	2.60	6	0	...	42.1201	48686.4225	
8	1	33.95	13	0	...	298.2529	371283.0489	
...	...	...	...	...	...	...	...	
200476	1	24.55	7	0	...	298.2529	570886.0249	
200478	1	3.78	7	0	...	28.3024	32648.8761	
200480	0	28.09	13	0	...	62.4100	233714.2336	
200486	1	18.39	6	0	...	183.0609	28815.0625	
200488	1	26.02	10	1	...	204.4900	265080.8196	

	f3	f4	f5	f6	f7	f8	f9	f10
0	1.368900e+10	688.5376	256	1.322704e+09	1747.24	625.0	0	1
2	4.225000e+09	486.2025	289	4.052572e+08	2840.89	729.0	9	1
4	1.853905e+09	163.5841	169	1.436882e+08	8500.84	676.0	0	1
6	2.916000e+09	6.7600	36	2.994278e+07	462.25	169.0	0	0
8	3.025000e+09	1152.6025	169	6.043731e+08	4872.04	1849.0	1	4
...	...	...	...	...	...	...	...	...
200476	4.624000e+09	602.7025	49	1.494905e+09	7072.81	289.0	16	4
200478	2.304000e+09	14.2884	49	5.346534e+07	1705.69	841.0	0	0
200480	1.764000e+09	789.0481	169	2.196028e+08	3203.56	729.0	0	0
200486	1.849000e+09	338.1921	36	5.290000e+06	1747.24	289.0	16	1
200488	1.936000e+09	677.0404	100	1.937107e+08	3757.69	529.0	0	4

[83135 rows x 32 columns]

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:7: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
import sys
```

In [183]: XS.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 83135 entries, 0 to 200488
Data columns (total 32 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   loan_amnt                            83135 non-null  float64
 1   int_rate                            83135 non-null  float64
 2   installment                          83135 non-null  float64
 3   mod_grade                           83135 non-null  int64
 4   sub_grade                           83135 non-null  object
 5   annual_inc                          83135 non-null  float64
 6   is_verified                         83135 non-null  int64
 7   dti                                 83135 non-null  float64
 8   open_acc                           83135 non-null  int64
 9   pub_rec                             83135 non-null  int64
10  revol_bal                           83135 non-null  float64
11  revol_util                          83135 non-null  float64
12  total_acc                          83135 non-null  float64
13  initial_list_status                 83135 non-null  int64
14  mort_acc                           83135 non-null  int64
15  pub_rec_bankruptcies                83135 non-null  int64
16  is_36mnths                         83135 non-null  int64
17  home_ownership_OTHERS               83135 non-null  int64
18  home_ownership_RENT                 83135 non-null  int64
19  purpose_debt_consolidation          83135 non-null  int64
20  purpose_others                      83135 non-null  int64
21  f0                                  83135 non-null  float64
22  f1                                  83135 non-null  float64
23  f2                                  83135 non-null  float64
24  f3                                  83135 non-null  float64
25  f4                                  83135 non-null  float64
26  f5                                  83135 non-null  int64
27  f6                                  83135 non-null  float64
28  f7                                  83135 non-null  float64
29  f8                                  83135 non-null  float64
30  f9                                  83135 non-null  int64
31  f10                                 83135 non-null  int64
dtypes: float64(16), int64(15), object(1)
memory usage: 20.9+ MB
```

In [184]: X\_train,X\_test,y\_train,y\_test=train\_test\_split(XS,y,test\_size=0.15,random\_state=2)

In [185]: X\_trainn,X\_dev,y\_trainn,y\_dev=train\_test\_split(X\_train,y\_train,test\_size=0.15,random\_state=2)

```
In [186]: print(X_trainn.shape)
          print(X_dev.shape)
          print(X_test.shape)
          print(y_trainn.shape)
          print(y_dev.shape)
          print(y_test.shape)
```

```
(60064, 32)
(10600, 32)
(12471, 32)
(60064,)
(10600,)
(12471,)
```

```
In [187]: X_trainn=pd.concat([X_trainn,y_trainn],axis=1)
```

```
In [188]: z=X_trainn.groupby('sub_grade')['y_obs'].mean().to_dict()
          X_trainn['sub_grade']=list(map(lambda x:z[x],X_trainn['sub_grade']))
          X_trainn['sub_grade']
```

```
Out[188]: 14488      0.192450
          94254      0.154802
          24670      0.341498
          20048      0.168598
          55286      0.075792
          ...
          71550      0.108300
          86534      0.347036
          117018     0.085448
          190226     0.213446
          184816     0.125066
          Name: sub_grade, Length: 60064, dtype: float64
```

```
In [189]: X_dev['sub_grade']=list(map(lambda x:z[x],X_dev['sub_grade']))
X_dev['sub_grade']
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

"""Entry point for launching an IPython kernel.

```
Out[189]: 39684      0.404795
16006      0.122754
165846     0.154802
151720     0.062593
15474      0.192450
...
192714     0.154802
38598      0.122754
132708     0.028833
90502      0.241294
179298     0.322265
Name: sub_grade, Length: 10600, dtype: float64
```

```
In [190]: X_test['sub_grade']=list(map(lambda x:z[x],X_test['sub_grade']))
X_test['sub_grade']
```

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

"""Entry point for launching an IPython kernel.

```
Out[190]: 56938      0.322265
63494      0.154802
120154     0.240930
26788      0.265314
77820      0.108300
...
143328     0.125066
177034     0.265314
48176      0.240930
49746      0.101064
110436     0.154802
Name: sub_grade, Length: 12471, dtype: float64
```

```
In [191]: cols=['loan_amnt','int_rate','installment','annual_inc','dti',
              'open_acc','revol_bal','revol_util','total_acc','mort_acc',
              'mod_grade']
d={}
for i in cols:
    d[i]=[X_trainn[i].mean(),X_trainn[i].std()]
print(d)
```

```
{'loan_amnt': [14491.336324587108, 8388.981997556977], 'int_rate': [13.78460392
2482301, 4.490693801881828], 'installment': [444.271407332167, 251.863757655670
16], 'annual_inc': [75394.66977274243, 48671.130821498744], 'dti': [17.68224127
597239, 8.108662503973125], 'open_acc': [11.59971030900373, 5.213968658647141],
'revol_bal': [16033.318560202451, 19607.84551585026], 'revol_util': [54.1856752
7970166, 23.89877657577768], 'total_acc': [25.88725359616409, 11.92963296750360
4], 'mort_acc': [1.8127164358018113, 2.148206675715537], 'mod_grade': [1.852124
400639318, 1.321668894368757]}
```




```
In [192]: cols=['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
               'open_acc', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
               'mod_grade']
for i in cols:
    mean=X_trainn[i].mean()
    std=X_trainn[i].std()
    X_trainn[i]=(X_trainn[i]-mean)/std
X_trainn.head(20)
```

Out[192]:

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	is_verified	d
14488	-0.869156	0.339679	-0.768596	0.111886	0.192450	-0.735439	0	0.55468
94254	0.418247	0.068006	0.681792	-0.644734	0.154802	0.053529	0	-0.05330
24670	0.060635	0.907520	-0.256136	0.868505	0.341498	-0.110839	1	0.43012
20048	-0.296977	-0.176945	-0.158822	0.111886	0.168598	-0.234116	0	0.85436
55286	1.610286	-1.310400	1.714652	-1.401353	0.075792	1.347931	1	-0.07057
138636	0.418247	0.442113	0.740554	0.868505	0.265314	-0.521760	1	0.63114
47806	-1.155246	-1.595434	-1.178778	-1.401353	0.040411	-0.850497	0	-1.59363
151938	-1.369813	-1.726816	-1.401398	-1.401353	0.028833	-0.727221	0	0.99372
153512	0.626854	-0.150223	0.882376	-0.644734	0.125066	-0.460122	1	-0.36901
151486	-1.012201	-1.582073	-1.031833	-1.401353	0.062593	0.505543	0	-0.58483
106528	2.444714	1.381835	1.917023	1.625124	0.391566	0.711003	1	0.50535
167006	-0.589027	0.713786	-0.412252	0.868505	0.298160	-0.932682	1	1.51415
129208	0.418247	0.994367	0.060622	0.868505	0.268987	-0.460122	1	0.56825
162968	-0.907898	1.604517	-0.735641	1.625124	0.391566	-1.035412	1	1.06401
163026	-0.726112	-0.319462	-0.650595	-0.644734	0.125066	0.074075	1	-0.53674
77624	1.610286	1.381835	1.180831	0.868505	0.322265	0.197352	1	0.20074
17694	-0.654589	0.402031	-0.514847	0.111886	0.241294	-0.809405	1	1.41302
20454	0.954665	-0.266018	1.228476	0.111886	0.192450	-0.521760	1	-1.61829
23510	-0.535385	-0.522103	-0.455768	-0.644734	0.125066	0.505543	1	-1.25695
161670	-0.535385	-0.176945	-0.426307	0.111886	0.168598	-0.829951	1	1.53635

20 rows × 33 columns



```
In [193]: cols=['loan_amnt','int_rate','installment','annual_inc','dti',
              'open_acc','revol_bal','revol_util','total_acc','mort_acc',
              'mod_grade']
for i in cols:
    X_dev[i]=(X_dev[i]-d[i][0])/d[i][1]
print(X_dev)
```

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	\
39684	2.444714	1.929634	2.109746	1.625124	0.404795	7.285743	
16006	-1.310211	-0.733651	-1.312302	-0.644734	0.122754	-0.398484	
165846	0.060635	-0.399627	0.213920	-0.644734	0.154802	-0.213570	
151720	1.610286	-1.372751	1.700318	-1.401353	0.062593	1.841037	
15474	-0.177773	-0.176945	-0.589769	0.111886	0.192450	0.197352	
...	...	...	...	...	...	...	
192714	-0.296977	-0.502061	-0.715353	-0.644734	0.154802	-0.727221	
38598	-0.714191	-0.399627	-0.643131	-0.644734	0.122754	-0.419030	
132708	-0.535385	-1.884921	-0.568249	-1.401353	0.028833	-0.012218	
90502	1.133471	0.667023	0.593013	0.111886	0.241294	0.300082	
179298	1.783132	0.842942	1.178012	0.868505	0.322265	-0.336846	

	is_verified	dti	open_acc	pub_rec	...	f1	f2	\
39684	1	-0.561405	0.268565	0	...	504.0025	951873.4096	
16006	1	-0.758725	0.652150	0	...	110.0401	12939.0625	
165846	0	0.622514	0.076773	0	...	143.7601	248153.4225	
151720	1	-0.996742	1.035735	0	...	58.0644	761291.1504	
15474	0	-0.737759	-1.073982	0	...	168.7401	87456.2329	
...	...	...	...	...	...	...	...	
192714	0	0.390664	-0.498605	0	...	132.9409	69748.8100	
38598	0	-0.192663	1.035735	0	...	143.7601	79687.6441	
132708	1	-1.187895	-0.690397	1	...	28.3024	90691.3225	
90502	1	1.224340	0.076773	0	...	281.5684	352396.5769	
179298	1	0.521388	-0.498605	0	...	308.7049	549036.5409	

	f3	f4	f5	f6	f7	f8	f9	f10
39684	1.849000e+11	172.3969	169	2.655959e+09	6756.84	1849.0	9	16
16006	3.136000e+09	132.9409	225	1.227101e+07	252.81	225.0	0	1
165846	4.225000e+09	516.6529	144	2.757260e+08	2237.29	1024.0	0	1
151720	2.722500e+10	92.1600	289	4.789970e+08	1672.81	1600.0	1	0
15474	7.225000e+09	136.8900	36	1.100191e+08	5227.29	324.0	1	4
...	...	...	...	...	...	...	...	...
192714	1.600000e+09	434.7225	81	1.029961e+09	4747.21	1296.0	9	1
38598	3.025000e+09	259.8544	289	9.329628e+07	1962.49	1024.0	1	1
132708	5.595040e+09	64.8025	64	1.651610e+07	2520.04	144.0	0	0
90502	8.100000e+09	762.3121	144	7.464917e+08	4692.25	900.0	4	4
179298	3.481000e+09	480.0481	81	3.013002e+08	3660.25	289.0	0	9

[10600 rows x 32 columns]

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))



```
In [194]: cols=['loan_amnt','int_rate','installment','annual_inc','dti',
              'open_acc','revol_bal','revol_util','total_acc','mort_acc',
              'mod_grade']
for i in cols:
    X_test[i]=(X_test[i]-d[i][0])/d[i][1]
print(X_test)
```

	loan_amnt	int_rate	installment	mod_grade	sub_grade	annual_inc	\
56938	1.875515	1.210369	2.648371	0.868505	0.322265	-0.069747	
63494	1.133471	-0.502061	1.379709	-0.644734	0.154802	-0.143230	
120154	0.179839	0.339679	-0.242279	0.111886	0.240930	-0.316300	
26788	-0.392340	0.406484	-0.691729	0.868505	0.265314	-0.710784	
77820	0.060635	-0.588908	0.189819	-0.644734	0.108300	0.094621	
...	...	...	...	...	...	...	
143328	-0.535385	-0.176945	-0.426307	-0.644734	0.125066	-0.521760	
177034	-0.296977	0.406484	-0.615140	0.868505	0.265314	-0.768313	
48176	-0.773793	0.045738	-0.678468	0.111886	0.240930	-0.747767	
49746	-0.535385	-0.844993	-0.482965	-0.644734	0.101064	1.121924	
110436	-0.010888	0.068006	0.192638	-0.644734	0.154802	-0.656337	

	is_verified	dti	open_acc	pub_rec	...	f1	f2
56938	1	0.951792	-0.882190	0	...	369.4084	1.234988e+06
63494	1	0.967824	1.227527	0	...	132.9409	6.268997e+05
120154	1	-1.043605	-0.690397	0	...	234.3961	1.468806e+05
26788	1	1.932225	0.268565	0	...	243.6721	7.292700e+04
77820	1	1.336566	-0.306812	0	...	124.0996	2.421427e+05
...	...	...	...	...	...	...	...
143328	1	0.267339	0.268565	0	...	168.7401	1.135016e+05
177034	1	-0.069338	-0.690397	1	...	243.6721	8.371764e+04
48176	1	1.917426	0.843942	0	...	195.7201	7.474209e+04
49746	0	-1.324786	-1.073982	0	...	99.8001	1.040901e+05
110436	1	2.103646	0.268565	1	...	198.5281	2.428420e+05

	f3	f4	f5	f6	f7	f8	f9	f10
56938	5.184000e+09	645.1600	49	3.330329e+09	7378.81	225.0	1	9
63494	4.681778e+09	651.7809	324	1.038565e+08	912.04	1296.0	1	1
120154	3.600000e+09	85.0084	64	2.815364e+07	5140.89	576.0	4	4
26788	1.664640e+09	1112.2225	169	2.984256e+08	4212.01	225.0	0	9
77820	6.400000e+09	813.3904	100	1.514377e+09	7921.00	676.0	36	1
...	...	...	...	...	...	...	...	...
143328	2.500000e+09	394.0225	169	3.053668e+07	6225.21	841.0	0	1
177034	1.444000e+09	293.0944	64	9.672100e+06	2034.01	841.0	4	9
48176	1.521000e+09	1104.2329	256	5.406661e+07	2016.01	400.0	0	4
49746	1.690000e+10	48.1636	36	1.314233e+08	3080.25	529.0	0	1
110436	1.887902e+09	1206.8676	169	3.661460e+07	7259.04	2116.0	81	1

[12471 rows x 32 columns]

C:\Users\Atul kumar\anaconda3\lib\site-packages\ipykernel\_launcher.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>


[table/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
In [195]: X_trainn=X_trainn.drop('y_obs',axis=1)
X_trainn.columns
```

```
Out[195]: Index(['loan_amnt', 'int_rate', 'installment', 'mod_grade', 'sub_grade',
               'annual_inc', 'is_verified', 'dti', 'open_acc', 'pub_rec', 'revol_bal',
               'revol_util', 'total_acc', 'initial_list_status', 'mort_acc',
               'pub_rec_bankruptcies', 'is_36mnths', 'home_ownership_OTHERS',
               'home_ownership_RENT', 'purpose_debt_consolidation', 'purpose_others',
               'f0', 'f1', 'f2', 'f3', 'f4', 'f5', 'f6', 'f7', 'f8', 'f9', 'f10'],
              dtype='object')
```

Making a LogisticRegression model

```
In [196]: train_cost=[]
val_cost=[]
lambdada=[]
for i in np.arange(0.01,500,2):
    lambdada.append(i)
    model=LogisticRegression(C=1/i,class_weight={0:0.1,1:0.6})
    model.fit(X_trainn,y_trainn)
    y_train_probs=model.predict_proba(X_trainn)
    y_val_probs=model.predict_proba(X_dev)
    train_loss=metrics.log_loss(y_trainn,y_train_probs)
    val_loss=metrics.log_loss(y_dev,y_val_probs)
    train_cost.append(train_loss)
    val_cost.append(val_loss)
plt.plot(lambdada,train_cost,label='train')
plt.plot(lambdada,val_cost,label='val')
plt.legend()
plt.show()
```



```
In [197]: z=np.argmin(val_cost)
z
```

```
Out[197]: 0
```

```
In [198]: model=LogisticRegression(C=1,class_weight={0:0.1,1:0.6})
model.fit(X_trainn,y_trainn)
print(model.score(X_trainn,y_trainn))
print(model.score(X_dev,y_dev))
print(model.score(X_test,y_test))
```

```
0.2590570058604156
0.25669811320754715
0.26124609093095985
```

```
In [199]: ypred=model.predict(X_test)
confusion_matrix(y_test,ypred)
```

```
Out[199]: array([[ 889, 9075],
 [ 138, 2369]], dtype=int64)
```

```
In [200]: print(precision_score(y_test,ypred))
print(recall_score(y_test,ypred))
```

```
0.20700803914715135
0.944954128440367
```

Removing the outliers which are on the wrong side of the plane  $\theta_T^T x + \theta_{not}$

```
In [201]: z=np.array(list(model.intercept_)+list(model.coef_[0]))
z=z.T
z
```

```
Out[201]: array([ 2.73464847e-18,  9.53859513e-19,  5.42256037e-18,  5.66917095e-19,
 5.71174031e-18,  1.17718627e-18, -1.77511311e-18,  2.65818229e-18,
 2.89722012e-18,  3.63015482e-19,  4.71306960e-19, -4.59448671e-19,
 1.59365616e-18, -7.35330550e-19,  1.11152473e-18, -1.78406958e-18,
 3.26092317e-19,  4.97488431e-19,  2.77981799e-19,  1.88506989e-18,
 1.91517594e-18,  5.73489833e-19,  9.90760756e-10,  1.29122727e-15,
 7.95629367e-13, -4.02482995e-12,  1.95358871e-15,  4.97865851e-16,
-1.74604483e-11,  1.33647727e-14,  1.74795612e-15,  5.97045872e-19,
 4.80592142e-17])
```

```
In [202]: X_trainn.insert(0,'x_not',1)
```

```
In [203]: z2=X_trainn.dot(z)
          z3=1/(1+np.exp(-z2))
          z3
```

```
Out[203]: 14488      0.510121
          94254      0.570642
          24670      0.549522
          20048      0.529288
          55286      0.660947
          ...
          71550      0.592478
          86534      0.566896
          117018     0.596300
          190226     0.492246
          184816     0.497507
          Length: 60064, dtype: float64
```

```
In [204]: X_trainn['log_loss']=-1*(y_trainn*np.log(z3)+(1-y_trainn)*(1-np.log(z3)))
```

```
In [205]: X_trainn['log_loss'].describe()
```

```
Out[205]: count      60064.000000
          mean        -1.165980
          std          0.884166
          min         -25.048947
          25%         -1.665072
          50%         -1.598013
          75%         -1.367070
          max           4.028935
          Name: log_loss, dtype: float64
```

```
In [206]: X_trainn=pd.concat([X_trainn,y_trainn],axis=1)
```

```
In [207]: q10=X_trainn['log_loss'].quantile(0.10)
          q90=X_trainn['log_loss'].quantile(0.90)
          X_trainn=X_trainn[(X_trainn['log_loss']>q10)&(X_trainn['log_loss']<q90)]
```

```
In [208]: X_trainn['log_loss'].describe()
```

```
Out[208]: count      48050.000000
          mean        -1.324091
          std          0.687728
          min         -1.688692
          25%         -1.655102
          50%         -1.598013
          75%         -1.457196
          max           0.607496
          Name: log_loss, dtype: float64
```


```
In [209]: y_trainn=X_trainn['y_obs']
y_trainn
```

```
Out[209]: 14488      0
          94254      0
          24670      0
          20048      0
          55286      0
          ..
          162398     0
          2676      0
          71550      0
          86534      1
          117018     0
Name: y_obs, Length: 48050, dtype: int64
```

```
In [210]: X_trainn=X_trainn.drop(['x_not','y_obs','log_loss'],axis=1)
```

Performing LogisticRegression after removing outliers

```
In [221]: train_cost=[]
val_cost=[]
lambdaa=[]
for i in np.arange(0.1,500,2):
    lambdaa.append(i)
    model=LogisticRegression(C=1/i,class_weight={0:0.1,1:1.2})
    model.fit(X_trainn,y_trainn)
    y_train_probs=model.predict_proba(X_trainn)
    y_val_probs=model.predict_proba(X_dev)
    train_loss=metrics.log_loss(y_trainn,y_train_probs)
    val_loss=metrics.log_loss(y_dev,y_val_probs)
    train_cost.append(train_loss)
    val_cost.append(val_loss)
plt.plot(lambdaa,train_cost,label='train')
plt.plot(lambdaa,val_cost,label='val')
plt.legend()
plt.show()
```





```
In [222]: z=np.argmin(val_cost)
          lbest=0.1+z*5
          lbest
```

Out[222]: 240.1

```
In [223]: model=LogisticRegression(C=1/lbest,class_weight={0:0.1,1:1.2})
          model.fit(X_trainn,y_trainn)
          print(model.score(X_trainn,y_trainn))
          print(model.score(X_dev,y_dev))
          print(model.score(X_test,y_test))
```

```
0.24276795005202914
0.36933962264150944
0.36612942025499157
```

```
In [224]: ypred=model.predict(X_test)
          confusion_matrix(y_test,ypred)
```

Out[224]: array([[2471, 7493],  
 [ 412, 2095]], dtype=int64)

```
In [225]: print(precision_score(y_test,ypred))
          print(recall_score(y_test,ypred))
```

```
0.21850229453483522
0.8356601515755884
```

EDA